

Analysis of song lyrics to match genre

Assignment 1: Basic text processing

Jernej Janež (63130077), Rok Marinšek (63130146), Luka Podgoršek (63130189)

November 11, 2018

1 NLP task

For our assignment we decided to analyze song lyrics, extract keywords that correspond to specific genres and try to classify song by its lyrics to corresponding genre. First we found a *dataset* that contained song lyrics, then we preprocessed the data, trained and tested a model and presented results with tables and graphs. Some similar solutions already exist, but perform similar task with neural networks or some other more complex methods. With this assignment we wanted to find out if our approach can provide satisfactory results by using simple natural language processing techniques.

2 Data

We searched the internet for appropriate datasets and found many different ones but in the end decided to use *380,000+ lyrics from MetroLyrics dataset* found on kaggle portal. This dataset had the attributes we needed to solve our task.

The dataset contained the following attributes:

- song title,
- year,
- artist
- genre,
- lyrics.

2.1 Data preparation

The data we found was stored in a *.csv* file. Because it contained more than *380 000* entries we decided to analyze songs that were released after year 2005 (latest songs in dataset). Afterwards we filtered the songs to match predefined genres which were *hip-hop*, *pop* and *metal*. Then we removed songs with lyrics that had less than 100 words and more than 1000 words which removed the outliers in the data.

When we finished preparing and selecting the data we focused on text preparation. First we removed special characters from text with regular expressions, converted words to lowercase and removed punctuations. Finally we removed non-english songs. This way we ended up with **64605** different songs.

| # | Genre | Number of different songs |
|---|---------|---------------------------|
| 0 | Hip-Hop | 19574 |
| 1 | Metal | 16187 |
| 2 | Pop | 28844 |

Table 1: Number of songs per genre

In the end we saved the filtered data into a *.csv* file and used it as input in our model class to train our model. You can also use this file to replicate our results.

3 Tools

To build our model we used a preprocessed file and logistic regression.

3.1 Train, test data and metrics

To train our model we used 80% of data and 20% to test our model. To measure score and performance of our model we used following metrics:

- accuracy,
- precision,
- recall,
- and f1 score.

In development phase we also played with regularization factor. We used above mentioned metrics to determine the best regularization factor. In the end we set it to 0.1.

3.2 Resources, tools and corpora

We used several different python libraries. Pandas was used for data structures and data purging. Nltk corpus was used to determine stopwords and for lemmatization. Langdetect library was used to remove non-english lyrics. To build our model we used sklearn and presented results with matplotlib.

4 Algorithm description

4.1 Data feature extraction

Features were determined by the Bag-Of-Words approach using the data of column *lyrics*. Our tokenizer has three roles: tokenize words using whitespaces (NLTK’s WhitespaceTokenizer), ignore stopwords and lemmatize each word. Each extracted feature has a corresponding TF-IDF value (using sklearn’s TfidfVectorizer).

4.2 Model

Our training data was constructed using the approach described in 4.1. To train our model for genre prediction we used sklearn’s multinomial LogisticRegression classifier.

To determine keywords for each genre we simply looked at the trained models weights and connected them to the corresponding feature. We decided to show only the 10 most significant words with the highest positive and negative weights, thinking those words represent the genre best.

5 Results

Before we could interpret our model and it’s predictions we had to test it. Testing phase consisted of 10 iterations. In each iteration we sampled the data with different random seeds, trained a model and tested it on sampled test data. In the end we averaged all metrics to score the performance of our model. You can find the results in the table bellow.

| Accuracy | Precision | Recall | F1 |
|----------|-----------|--------|-------|
| 0.823 | 0.826 | 0.823 | 0.823 |

Table 2: Model scores

Beside performance scores our model returned a list of keywords that have the most and the least value for classification. These results are presented in the tables bellow. Graphical presentation of tables is included in the appendix graphs section.

| (+) Keywords | Weight | (-) Keywords | Weight | (+) Keywords | Weight | (-) Keywords | Weight |
|--------------|--------|--------------|--------|--------------|--------|--------------|--------|
| death | 1.97 | love | -2.97 | nigga | 3.71 | i’ve | -1.20 |
| blood | 1.88 | baby | -2.31 | shit | 2.71 | eye | -1.09 |
| dead | 1.85 | like | -2.29 | yo | 2.48 | lie | -1.02 |
| fucking | 1.65 | oh | -2.13 | ain’t | 2.46 | dream | -0.94 |
| pain | 1.57 | got | -1.98 | like | 2.38 | alone | -0.77 |
| die | 1.56 | girl | -1.96 | bitch | 2.25 | inside | -0.77 |
| end | 1.53 | i’m | -1.95 | girl | 1.76 | heart | -0.77 |
| lie | 1.46 | ain’t | -1.81 | ya | 1.75 | end | -0.76 |
| hate | 1.43 | yeah | -1.81 | cause | 1.69 | there’s | -0.75 |
| fear | 1.43 | get | -1.64 | got | 1.67 | away | -0.75 |

Table 3: Keywords for: **metal** (left) & **hip-hop** (right)

| (+) Keywords | Weight | (-) Keywords | Weight |
|--------------|--------|--------------|--------|
| love | 2.29 | nigga | -2.45 |
| oh | 1.93 | shit | -2.41 |
| heart | 1.56 | fuck | -2.14 |
| gonna | 1.29 | death | -1.76 |
| boy | 1.22 | fucking | -1.66 |
| baby | 1.16 | dead | -1.61 |
| 'cause | 1.10 | die | -1.57 |
| could | 1.01 | bitch | -1.55 |
| there's | 0.97 | blood | -1.43 |
| kiss | 0.97 | hate | -1.37 |

Table 4: Keywords for **pop**

5.1 POS tagging

To gain further insight into our model's keywords we decided to POS tag them. We extracted the 10 most (+) significant nouns, adjectives and verbs for each genre. The resulting words are visible in the tables 5, 6 and 7 below. We find the results not to be surprising, especially the nouns are something that should be recognizable to each fan of the genres (and also confirm certain stereotypes we have about the genres).

| Top verbs | Weight | Top nouns | Weight | Top adjectives | Weight |
|-----------|--------|-----------|--------|----------------|--------|
| got | 1.67 | nigga | 3.71 | whole | 0.91 |
| get | 1.67 | shit | 2.71 | big | 0.82 |
| keep | 0.68 | yo | 2.48 | real | 0.79 |
| go | 0.62 | ain't | 2.46 | black | 0.60 |
| hundred | 0.57 | bitch | 2.25 | fresh | 0.49 |
| said | 0.54 | girl | 1.76 | high | 0.48 |
| know | 0.54 | ya | 1.75 | low | 0.48 |
| started | 0.45 | cause | 1.69 | nice | 0.45 |
| came | 0.44 | 'em | 1.66 | next | 0.42 |
| slow | 0.41 | i'm | 1.66 | new | 0.41 |

Table 5: Hip-hop POS tagging

5.2 Understanding wrong classifications

If you look at our confusion matrix you can notice that most commonly miss-classified songs are pop songs. We didn't expect to classify every song correctly, but pop songs have larger error rate than metal or hip-hop songs. We tried to understand why this was happening and our hypothesis was that many pop song artists feature artists from hip-hop or metal and therefore correlate with actual hip-hop and metal songs.

To test this we extracted songs that were not correctly classified. Afterwards we checked if song titles contained words *"featuring"* or *"ft"* or *"feat"*. We counted the percentage of these songs and you can see results in the table below.

| Top verbs | Weight | Top nouns | Weight | Top adjectives | Weight |
|-----------|--------|-----------|--------|----------------|--------|
| fucking | 1.65 | death | 1.97 | dead | 1.85 |
| dying | 0.83 | blood | 1.88 | final | 1.14 |
| hatred | 0.67 | pain | 1.57 | eternal | 0.91 |
| lost | 0.66 | die | 1.56 | alive | 0.69 |
| forced | 0.60 | end | 1.53 | black | 0.61 |
| suffering | 0.58 | lie | 1.46 | empty | 0.54 |
| killing | 0.57 | hate | 1.43 | unknown | 0.51 |
| damned | 0.56 | fear | 1.43 | false | 0.50 |
| wing | 0.51 | soul | 1.36 | free | 0.49 |
| screaming | 0.50 | hell | 1.30 | last | 0.48 |

Table 6: Metal POS tagging

| Top verbs | Weight | Top nouns | Weight | Top adjectives | Weight |
|-----------|--------|-----------|--------|----------------|--------|
| know | 0.90 | love | 2.29 | little | 0.76 |
| say | 0.80 | oh | 1.93 | easy | 0.53 |
| loving | 0.73 | heart | 1.56 | happy | 0.51 |
| darling | 0.68 | gonna | 1.29 | good | 0.45 |
| find | 0.63 | boy | 1.22 | small | 0.39 |
| feeling | 0.61 | baby | 1.16 | lucky | 0.32 |
| believe | 0.59 | 'cause | 1.10 | much | 0.28 |
| go | 0.58 | there's | 0.97 | tiny | 0.26 |
| let | 0.56 | kiss | 0.97 | old | 0.24 |
| dancing | 0.55 | ooh | 0.96 | impossible | 0.23 |

Table 7: Pop POS tagging

6 Conclusion

7 Github repository

Github repository: <https://github.com/marok39/onj-02>

Appendix

A Graphs

Here you can find the visualization of all results. These files can be found alongside the report in */img* directory.




Figure 1: Confusion matrix




Figure 2: Keywords visualization for pop

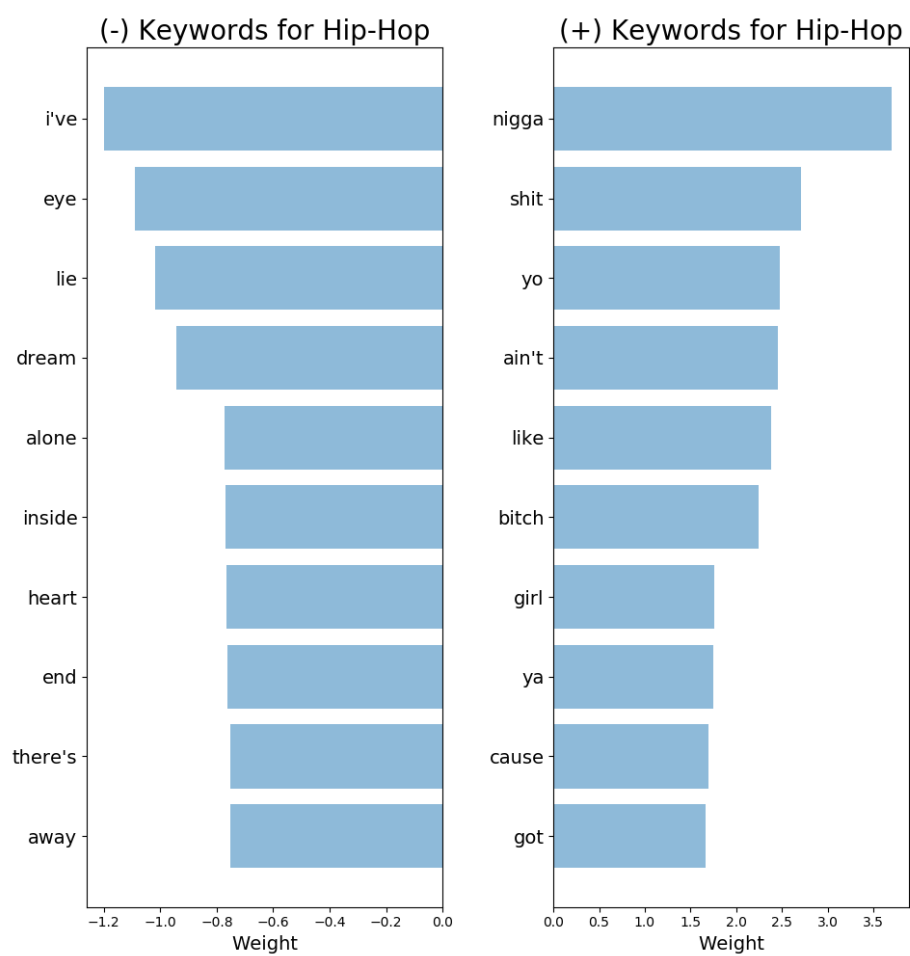


Figure 3: Keywords visualization for hip-hop




Figure 4: Keywords visualization for metal